# Range Query Processing for Monitoring Applications over Untrustworthy Clouds

Hoang Van Tran, Tristan Allard, Laurent D'Orazio, and Amr El Abbadi

{van-hoang.tran, laurent.dorazio}@univ-rennes1.fr,
tristan.allard@irisa.fr, amr@ucsb.edu

## Introduction

- Use case: Real-time monitoring of the spatial spread of seasonal epidemics
  - Outsourced data in untrusted clouds
  - Range queries
  - High rate data arrivals
- Challenge: Performance problems (e.g., bottlenecks and overload)
- Proposal: PINED-RQ++, an extension of PINED-RQ [2]

## PINED-RQ [2]

**Objective:** Efficient range query processing in untrusted clouds via secure and low-storage consumption indexes

**Privacy model:** Probabilistic relaxation of a computational variant of differential privacy

**PINED-RQ index:** Built in 2 steps

**Step 1: Construction**

- An index is constructed based on a B+Tree
- The set of all nodes is defined as a histogram covering the domain of an indexed attribute

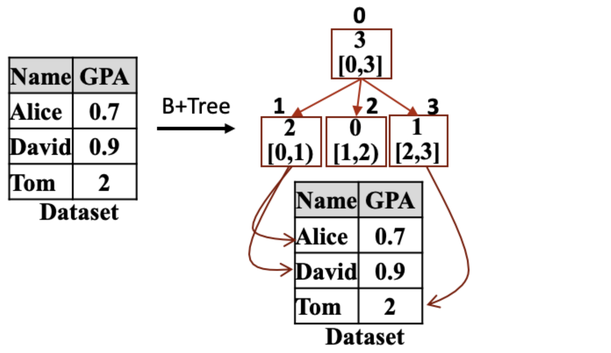


**Step 2: Perturbation**

- Laplace noise [1] are added to all counts
- The (positive/negative) noise makes the counts changing
- To keep the consistency between a leaf node's noisy count and the number of pointers it holds:
  - Positive noise: Add dummy tuples to the dataset
  - Negative noise: Remove real tuples from the dataset
  - Conceal removed tuples in overflow arrays

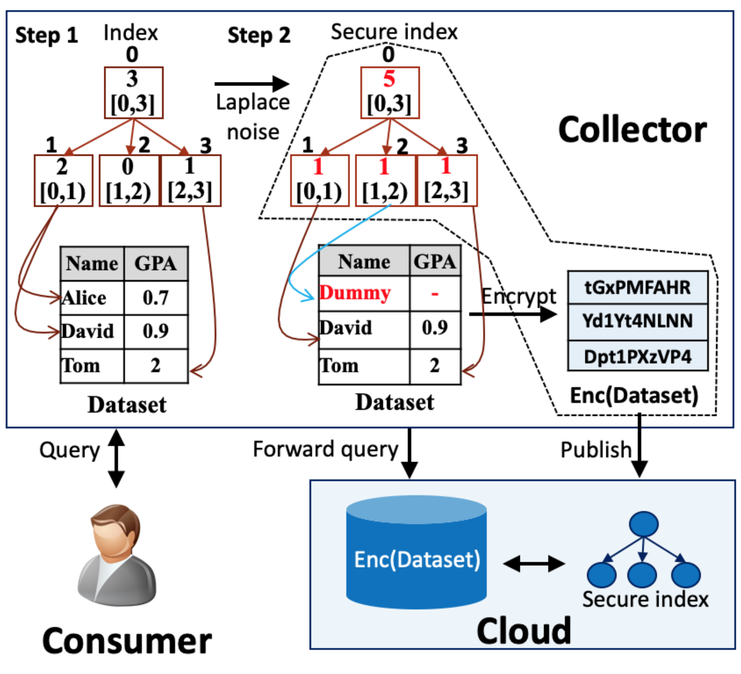


### Limitations

- Publish data in batches
- Partially process queries at the collector

  ⇒ Potential bottlenecks

## PINED-RQ++

**Intuition:** Send a new tuple to the cloud immediately when it arrives at the collector

**Proposal:** Index template

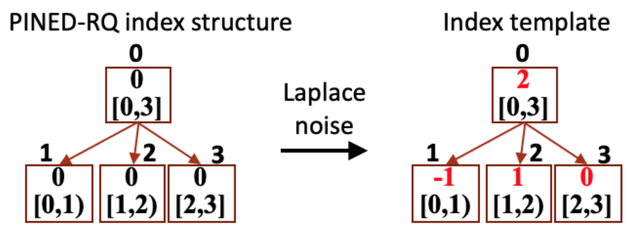


### Index template updating

- Initiate an index template
- Update the index template as a new tuple arrives
- Encrypt and forward the tuple to the cloud
- Associate the index template with unindexed data to produce a secure index
- Use unique random numbers and a matching table to keep the relationship between published tuples and its index

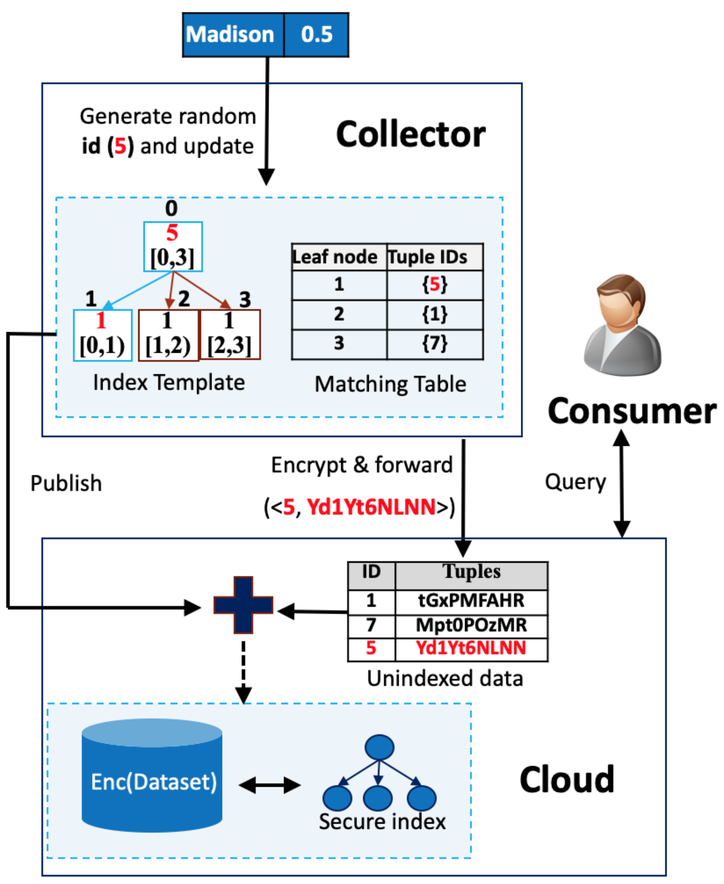


### Challenges

How to manage dummy and removed tuples?

- Dummy tuples: The collector sends dummy data according to the actual distribution of the sending time of the real tuples
- Removed tuples: If a leaf node receives negative noise $c$, move its first $c$ tuples to the corresponding overflow array

### Query processing

A query is only processed at the cloud holding indexed and unindexed data at time

### Parallel PINED-RQ++

- To enhance the collector's throughput
- Multiple clones of the original template

**Assumption:** An individual has a single tuple

## Experimental Results

### Benchmark Environment

The proposal was evaluated on NASA log [3] and Gowalla [4]

### Results

- Network traffic (Maximum data rate)

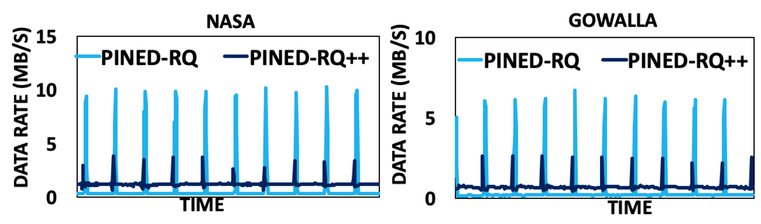


- Publishing time

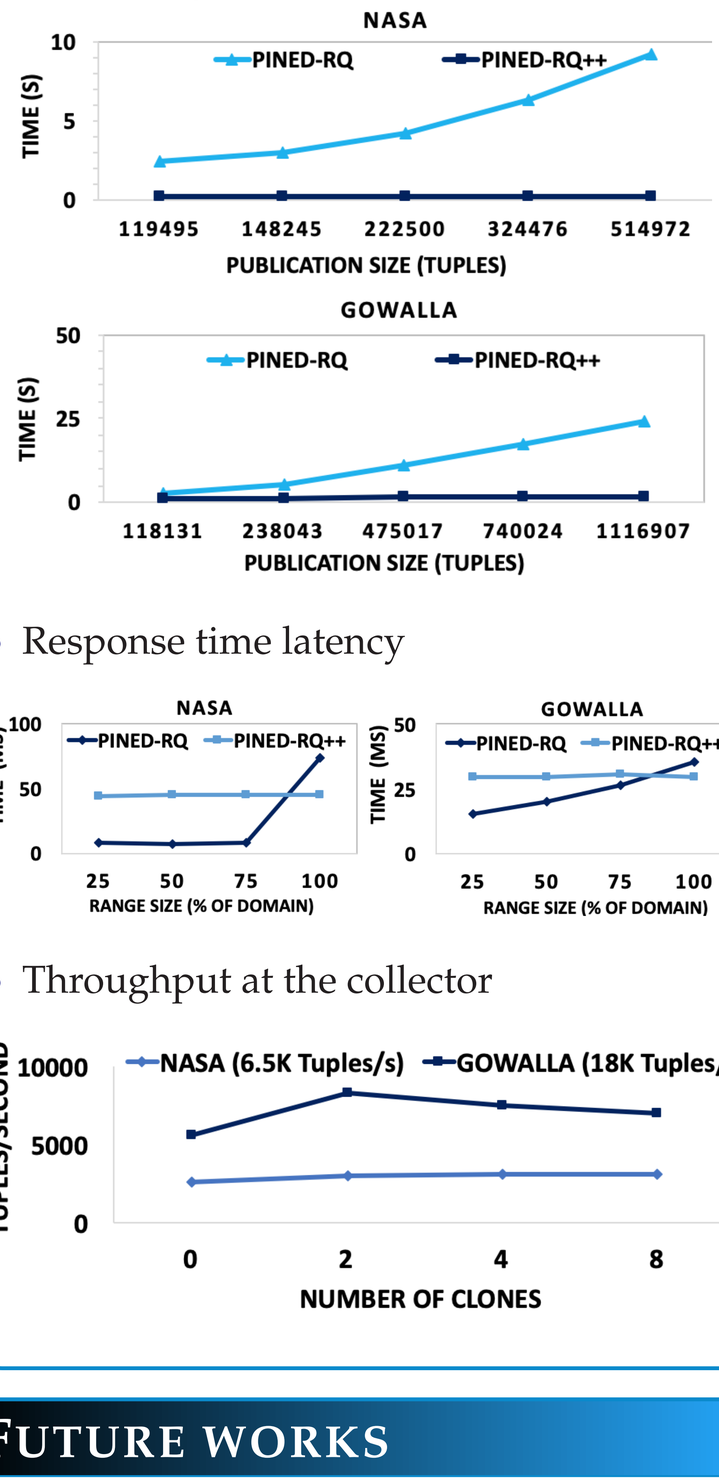


- Response time latency




- Throughput at the collector




## Future Works

- Query performance improvement (caching)
- Dynamic adaptation to query workload

## References

[1] C. Dwork, F. McSherry, K. Nissim, and A. Smith Calibrating Noise to Sensitivity in Private Data Analysis. In *TCC '06*. Pages 265–284.

[2] C. Sahin, T. Allard, R. Akbarina, A. El Abbadi, and E. Pacitti. A Differentially Private Index for Range Query Processing in Clouds. In *ICDE'18*. Pages 857–868.

[3] NASA Log (1996). http://ita.ee.lbl.gov/html/contrib/NASA-HTTP.html

[4] J. Leskovec and A. Krevl. SNAP Datasets: Stanford Large Network Dataset Collection. June, 2014. http://snap.stanford.edu/data.

## Acknowledgement

This work was partially supported by ARED Pôle d'Excellence Cyber/Lannion-Trégor Communauté.